

# HEURISTIC SEARCH

Ivan Bratko

Faculty of Computer and Information Sc.

University of Ljubljana

# Best-first search

- Best-first: most usual form of heuristic search
- Evaluate nodes generated during search
- Evaluation function  
 $f: \text{Nodes} \rightarrow \mathbb{R}^+$
- Convention: lower  $f$ , more promising node;  
higher  $f$  indicates a more difficult problem
- Search in directions where  $f$ -values are lower

# Heuristic search algorithms

- A\*, perhaps the best known algorithm in AI
- Hill climbing, steepest descent
- Beam search
- IDA\*
- RBFS

# Heuristic evaluation in A\*

*The question: How to find successful f?*


Algorithm A\*:

$f(n)$  estimates cost of best solution through  $n$


$$f(n) = g(n) + h(n)$$

known


heuristic guess,  
estimate of path cost  
from  $n$  to nearest goal




## Route search in a map




# Best-first search for shortest route



# Expanding tree within Bound



## g-value and f-value of a node



$$g(M) = g(N) + C$$

$$f(M) = g(M) + h(M)$$



# Admissibility of A\*

- A search algorithm is *admissible* if it is guaranteed to always find an optimal solution
- Is there any such guarantee for A\*?
- Admissibility theorem (Hart, Nilsson, Raphael 1968):  
A\* is admissible if  $h(n) \leq h^*(n)$  for all nodes  $n$  in state space.  $h^*(n)$  is the actual cost of min. cost path from  $n$  to a goal node.

# Admissible heuristics

- $A^*$  is admissible if it uses an optimistic heuristic estimate
- Consider  $h(n) = 0$  for all  $n$ .  
This trivially admissible!
- However, what is the drawback of  $h = 0$ ?
- How well does it guide the search?
- Ideally:  $h(n) = h^*(n)$
- Main difficulty in practice: Finding heuristic functions that guide search well and are admissible

# Finding good heuristic functions

- Requires problem-specific knowledge
- Consider 8-puzzle
- $h = \text{total\_dist} = \text{sum of Manhattan distances of tiles from their "home" squares}$

1	3	4
8	5	
7	6	2

$$\text{total\_dist} = 0+3+1+1+2+0+0+0=7$$

- Is  $\text{total\_dist}$  admissible?

# Heuristics for 8-puzzle

*sequence\_score* : assess the order of tiles;  
count 1 for tile in middle, and 2 for each tile on edge not followed by proper successor clockwise

1	3	4
8	5	
7	6	2

$$\text{sequence\_score} = 1 + 2 + 0 + 2 + 2 + 0 + 0 + 0 = 7$$

$$h = \text{total\_dist} + 3 * \text{sequence\_score}$$

$$h = 7 + 3 * 7 = 28$$

Is this heuristic function admissible?

# Three 8-puzzles

1	3	4
8		2
7	6	5

(a)

2	8	3
1	6	4
7		5

(b)




2	1	6
4		8
7	5	3

(c)

Puzzle c requires 18 steps

$A^*$  with this  $h$  solves (c) almost without any branching

# A task-scheduling problem and two schedules



Optimal

# Heuristic function for scheduling

- Fill tasks into schedule from left to right
- A heuristic function:

For current, partial schedule:

$F_J$  = finishing time of processor J current engagement

$Fin = \max F_J$

$Finall = ( \text{SUM}_W(D_W) + \text{SUM}_J(F_J) ) / m$

W: waiting tasks;  $D_W$ : duration of waiting task W

$h = \max( Finall - Fin, 0 )$

- Is this heuristic admissible?

# Space Saving Techniques for Best-First Search


- Space complexity of  $A^*$ : depends on  $h$ , but it is roughly  $b^d$
- Space may be critical resource
- Idea: trade time for space, similar to iterative deepening
- Space-saving versions of  $A^*$ :
- IDA\* (Iterative Deepening  $A^*$ )
- RBFS (Recursive Best First Search)



# IDA\*, Iterative Deepening A\*

- Introduced by Korf (1985)
- Analogous idea to iterative deepening
- Iterative deepening:
  - Repeat depth-first search within increasing depth limit
- IDA\*:
  - Repeat depth-first search within increasing f-limit

## f-values form relief



IDA\*: Repeat depth-first within f-limit increasing:  
 $f_0, f_1, f_2, f_3, \dots$

# IDA\*

Bound := f(StartNode);

SolutionFound := false;

## **Repeat**

perform depth-first search from StartNode, so that

a node N is expanded only if  $f(N) \leq \text{Bound}$ ;

**if** this depth-first search encounters a goal node with  $f \leq \text{Bound}$

**then** SolutionFound = true

## **else**

compute new bound as:

Bound = min {  $f(N)$  | N generated by this search,  $f(N) > \text{Bound}$  }


**until** solution found.

# IDA\* performance

- Experimental results with 15-puzzle good  
( $h$  = Manhattan distance, many nodes same  $f$ )
- However: May become very inefficient when nodes do not tend to share  $f$ -values (e.g. small random perturbations of Manhattan distance)

# IDA\*: problem with non-monotonic f

- Function  $f$  is *monotonic* if:  
for all nodes  $n_1, n_2$ : if  $s(n_1, n_2)$  then  $f(n_1) \leq f(n_2)$
- Theorem: If  $f$  is monotonic then IDA\* expands nodes in best-first order
- Example with non-monotonic  $f$ :




$f$ -limit = 5, “3” expanded  
before “1”, “2”

# Linear Space Best-First Search

- Linear Space Best-First Search
- RBFS (Recursive Best First Search), Korf 93
- Similar to  $A^*$  implementation in Bratko (86; 2001), but saves space by iterative re-generation of parts of search tree

# Node values in RBFS



$f(N)$  = 'static' f-value

$F(N)$  = backed-up f-value,

i.e. currently known lower bound  
on cost of solution path through  $N$

$F(N) = f(N)$  if  $N$  has (never) been expanded

$F(N) = \min_i F(N_i)$  if  $N$  has been expanded




# Simple Recursive Best-First Search

## SRBFS

- First consider SRBFS, a simplified version of RBFS
- Idea: Keep expanding subtree within F-bound determined by siblings
- Update node's F-value according to searched subtree
- SRBFS expands nodes in best-first order




# Example: Searching tree with non-monotonic f-function



# SRBFS can be very inefficient


Example: search tree with  
 $f(\text{node}) = \text{depth of node}$



Parts of tree repeatedly re-explored;  
 $f$ -bound increases in unnecessarily small steps


# RBFS, improvement of SRBFS

- F-values not only backed-up from subtrees, but also *inherited* from parents
- Example: searching tree with  $f = \text{depth}$
- At some point, F-values are:




In RBFS, “7” is expanded, and F-value inherited:

SRBFS











RBFS



## F-values of nodes

- To save space RBFS often removes already generated nodes
- But: If  $N_1, N_2, \dots$  are deleted, essential info. is retained as  $F(N)$
- Note: If  $f(N) < F(N)$  then (part of)  $N$ 's subtree must have been searched

# Searching the map with RBFS



# Algorithm RBFS

RBFS ( N, F(N) , Bound)

if  $F(N) > \text{Bound}$  then return  $F(N)$  ;

if goal(N) then exit search;

if N has no children then return  $\infty$ ;

for each child  $N_i$  of N do

    if  $f(N) < F(N)$  then  $F(N_i) := \max(F(N), f(N_i))$

        else  $F(N_i) := f(N_i)$  ;

Sort  $N_i$  in increasing order of  $F(N_i)$  ;

if only one child then  $F(N_2) := \infty$ ;

while  $F(N_1) \leq \text{Bound}$  and  $F(N_1) < \infty$  do

$F(N_1) := \text{RBFS}(N_1, F(N_1), \min(\text{Bound}, F(N_2)))$

    insert  $N_1$  in sorted order

return  $F(N_1)$  .

# Properties of RBFS

- RBFS( Start,  $f(\text{Start})$ ,  $\infty$ ) performs complete best-first search
- Space complexity:  $O(bd)$   
(linear space best-first search)
- RBFS explores nodes in best-first order even with non-monotonic  $f$   
That is: RBFS expands “*open*” nodes in best-first order




# Summary of concepts in best-first search

- Evaluation function:  $f$
- Special case ( $A^*$ ):  $f(n) = g(n) = h(n)$
- $h(n)$  admissible if  $h(n) \leq h^*(n)$
- Sometimes in literature:  $A^*$  defined with admissible  $h$
- Algorithm “respects” best-first order if already generated nodes are expanded in best-first order according to  $f$
- $f$  is defined with intention to reflect the “goodness” of nodes; therefore it is desirable that an algorithm respects best-first order
- $f$  is monotonic if for all  $n_1, n_2$ :  $s(n_1, n_2) \Rightarrow f(n_1) \leq f(n_2)$

# Best-first order

- Algorithm “respects” best-first order if generated nodes are expanded in best-first order according to  $f$
- $f$  is defined with intention to reflect the “goodness” of nodes; therefore it is desirable that an algorithm respects best-first order
- $f$  is monotonic if for all  $n_1, n_2$ :  $s(n_1, n_2) \Rightarrow f(n_1) \leq f(n_2)$
- $A^*$  and RBFS respect best-first order
- IDA\* respects best-first order if  $f$  is monotonic

Problem. How many nodes are generated by A\*, IDA\* and RBFS? Count also re-generated nodes




# SOME OTHER BEST-FIRST TECHNIQUES

- Hill climbing, steepest descent, greedy search: special case of  $A^*$  when the successor with best  $F$  is retained only; no backtracking
- Beam search: special case of  $A^*$  where only some limited number of open nodes are retained, say  $W$  best evaluated open nodes ( $W$  is beam width)
- In some versions of beam search,  $W$  may change with depth. Or, the limitation refers to number of successors of an expanded node retained

# REAL-TIME BEST FIRST SEARCH

- With limitation on problem-solving time, agent (robot) has to make decision before complete solution is found
- RTA\*, real-time A\* (Korf 1990):
  - Agent moves from state to next best-looking successor state after fixed depth lookahead
  - Successors of current node are evaluated by backing up f-values from nodes on lookahead-depth horizon
  - $f = g + h$

# RTA\* planning and execution




## RTA\* and alpha-pruning

- If  $f$  is monotonic, *alpha-pruning* is possible (with analogy to alpha-beta pruning in minimax search in games)
- Alpha-pruning: let  $\text{best\_f}$  be the best  $f$ -value at lookahead horizon found so far, and  $n$  be a node encountered before lookahead horizon; if  $f(n) \geq \text{best\_f}$  then subtree of  $n$  can be pruned
- Note: In practice, many heuristics are monotonic (Manhattan distance, Euclidean distance)

# Alpha pruning

- Prune  $n$ 's subtree if  $f(n) \geq \alpha$
- If  $f$  is monotonic then all descendants of  $n$  have  $f \geq \alpha$





## RTA\*, details

- Problem solving = planning stage + execution stage
- In RTA\*, planning and execution interleaved
- Distinguished states: start state, current state
- Current state is understood as actual physical state of robot reached after physically executing a move
- Plan in current state by fixed depth lookahead, execute one move to reach next current state

## RTA\*, details

- $g(n)$ ,  $f(n)$  are measured relative to current state (not start state)
- $f(n) = g(n) + h(n)$ , where  $g(n)$  is cost from current state to  $n$  (not from start state)

## RTA\* main loop, roughly

- current state  $s := \text{start\_state}$
- While goal not found:
  - Plan: evaluate successors of  $s$  by fixed depth lookahead;
  - $\text{best\_s} :=$  successor with min. backed-up  $f$
  - $\text{second\_best\_f} := f$  of second best successor
  - Store  $s$  among “visited nodes” and store
$$f(s) := f(\text{second\_best\_f}) + \text{cost}(s, \text{best\_s})$$
  - Execute: current state  $s := \text{best\_s}$

## RTA\*, visited nodes

- Visited nodes are nodes that have been (physically) visited (i.e. robot has moved to these states in past)
- Idea behind storing  $f$  of a visited node  $s$  as:  
$$f(s) := f(\text{second\_best\_f}) + \text{cost}(s, \text{best\_s})$$
- If  $\text{best\_s}$  subsequently turns out as bad, problem solver will return to  $s$  and this time consider  $s$ 's second best successor;  $\text{cost}(s, \text{best\_s})$  is added to reflect the fact that problem solver had to pay the cost of moving from  $s$  to  $\text{best\_s}$ , in addition to later moving from  $\text{best\_s}$  to  $s$

## RTA\* lookahead

For node  $n$  encountered by lookahead:

- if  $\text{goal}(n)$  then return  $h(n) = 0$ ,  
don't search beyond  $n$
- if  $\text{visited}(n)$  then return  $h(n) = \text{stored } f(n)$ ,  
don't search beyond  $n$
- if  $n$  at lookahead horizon then evaluate  $n$  statically  
by heuristic function  $h(n)$
- if  $n$  not at lookahead horizon then generate  $n$ 's  
successors and back up  $f$  value from them

# LRTA\*, Learning RTA\*

- Useful when successively solving multiple problem instance with the same goal
- Trial = Solving one problem instance
- Save table of visited nodes with their backed-up h values
- Note: In table used by LRTA\*, store the best successors' f (rather than second best f as in RTA\*). Best f is appropriate info. to be transferred *between* trials, second best is appropriate *within* a single trial

# In RTA\*, pessimistic heuristics better than optimistic (Sadikov, Bratko 2006)

- Traditionally, optimistic heuristics preferred in A\* search (admissibility)
- Surprisingly, in RTA\* pessimistic heuristics perform better than optimistic (solutions closer to optimal, less search, no pathology)